

# ACC Point Spread Predictions and Prediction Intervals (2026 Season)

Team: YOUR\_TEAM\_NAME | Team members: YOUR\_NAME(S)

## 1) Problem and target

Goal: predict the game point spread (home score – away score) for 78 ACC men's games, and provide prediction intervals ( $ci_{lb}$ ,  $ci_{ub}$ ) with at least 70% empirical coverage.

## 2) Data construction

Historical data were assembled from prior ACC seasons (2023–2025) plus any completed games in the current season (2026) using an ESPN-based scrape (CBBpy). For each played game, we compute the observed point spread and a feature vector built from team-level efficiency ratings and context flags.

## 3) Model and training procedure (point prediction)

Primary point-spread model: a tuned Gradient Boosting regression model (scikit-learn GradientBoostingRegressor, loss='absolute\_error'). Hyperparameters were tuned via RandomizedSearchCV using K-fold cross-validation on the training split only (to avoid leaking test information).

Two evaluation schemes were run:

- Season holdout (train 2023–2025, test 2026): best MAE  $\approx 8.36$  (Lasso/linear models).
- Random split (mixed 2023–2026): best MAE  $\approx 7.89$  (GradientBoosting\_TUNED).

For the submitted point predictions, we use the random-split tuned Gradient Boosting model because it leverages the largest training set including current-season games and performed best under the mixed-year test.

## 4) Predictors (features)

Each upcoming game is represented by home-minus-away differences plus context flags:

- net\_rating\_diff ( $adj\_off\_eff - adj\_def\_eff$  differential)
- efg\_pct\_diff
- luck\_diff
- adj\_off\_eff\_diff
- home\_advantage (team-level home/away margin differential)
- rating\_luck\_interaction =  $net\_rating\_diff \times luck\_diff$
- shooting\_strength\_interaction =  $efg\_pct\_diff \times adj\_off\_eff\_diff$
- is\_conference (1/0)
- is\_neutral (1/0)

## 5) Prediction intervals (method)

Intervals are constructed using an empirical residual calibration approach:

- 1) On a validation set (held out from training), compute residuals  $r = y - \hat{y}$ .
- 2) Let  $q$  be the smallest half-width such that the empirical coverage of  $[\hat{y} - q, \hat{y} + q]$  is at least 70%. Operationally,  $q$  is chosen as an empirical quantile of  $|r|$  that meets the 70% coverage constraint.
- 3) For each new game, report:  $ci_{lb} = \hat{y} - q$  and  $ci_{ub} = \hat{y} + q$ .

This method directly targets the competition rule ( $\geq 70\%$  coverage) while keeping intervals as tight as possible (small average width).

## 6) Model assessment and selection

Point predictions were evaluated using MAE and  $R^2$  on out-of-sample splits. Prediction intervals were assessed via empirical coverage (must be  $\geq 70\%$ ) and average interval width. The submitted interval half-width  $q$  was selected to satisfy the coverage constraint on validation data.